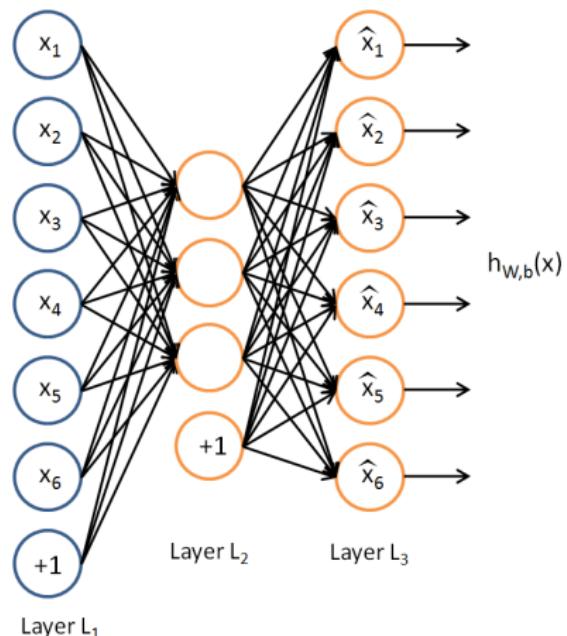


## More Neural Networks

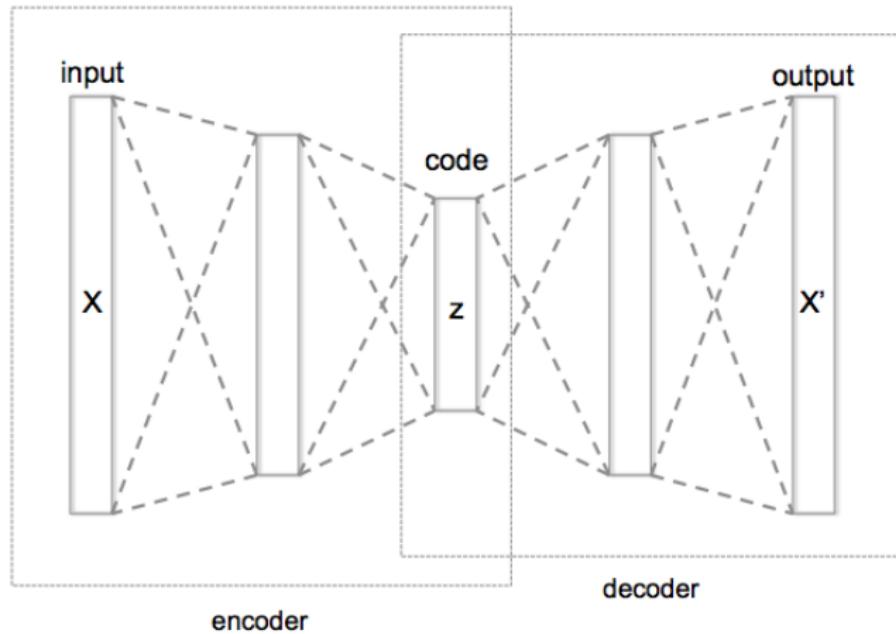
Mladen Kolar ([mkolar@chicagobooth.edu](mailto:mkolar@chicagobooth.edu))

# Autoencoder



Network trained to reproduce its input at the output layer.  
Usually tie the weights that go into and out of the hidden layer.

# Autoencoder



# Autoencoder

## Loss function

- ▶ For real valued inputs, try to find weights such that

$$\frac{1}{2} \sum_k (x_k - \hat{x}_k)^2$$

is minimized

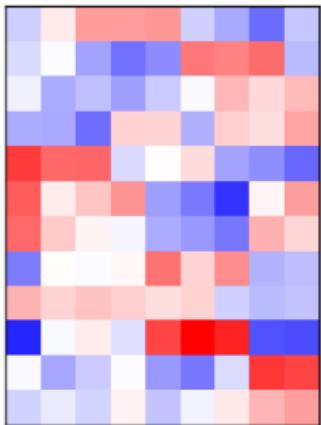
- ▶ For binary input cross entropy is used, which is similar to deviance

## Fitting autoencoder

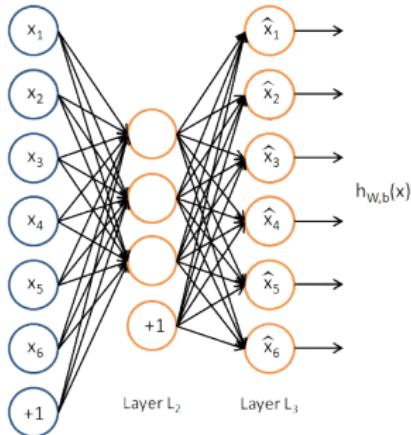
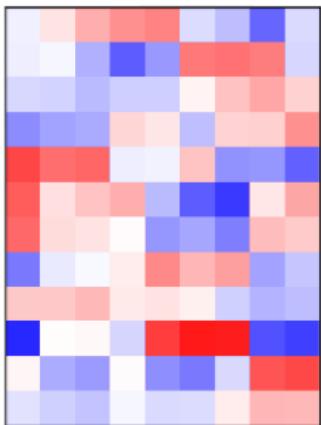
- ▶ Same tricks as before
- ▶ Greedy learning of stacked autoencoders
- ▶ <https://www.cs.toronto.edu/~hinton/science.pdf>

# Autoencoder

Original Data

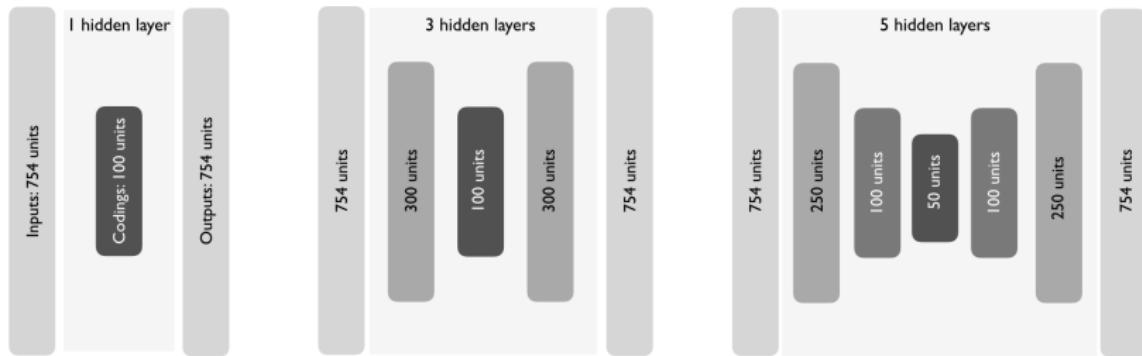


Reconstruction



$$\min \sum_{i=1}^n \|x_i - \sigma(W'^\top \sigma(Wx_i + b) + b')\|_2^2$$

# Stacked (or deep) autoencoders

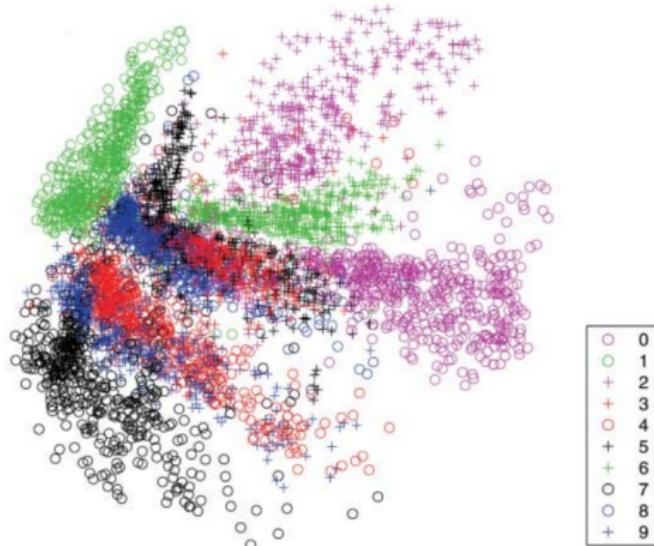


Additional depth can help

- ▶ codes represent more complex, nonlinear relationships
- ▶ reduced computational cost
- ▶ better data compression than shallower, or linear autoencoders.

# Autoencoder: Why are they useful?

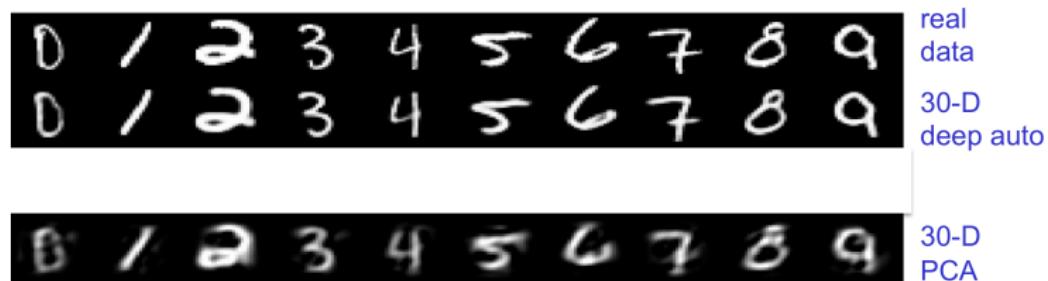
Learning compressed representation of the input distribution  
(dimensionality reduction)



Autoencoder structure: 784 — 1000 — 500 — 250 — 2

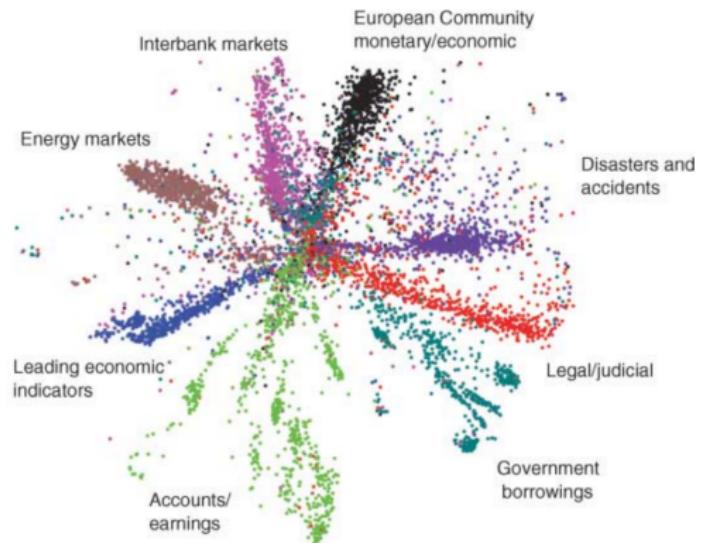
<https://www.cs.toronto.edu/~hinton/science.pdf>

A comparison of methods for compressing digit images to 30 real numbers



# Autoencoder: Why are they useful?

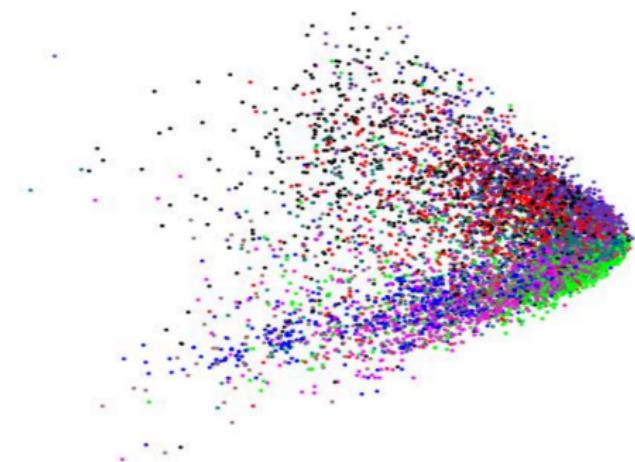
Information retrieval: 804,414 newswire stories



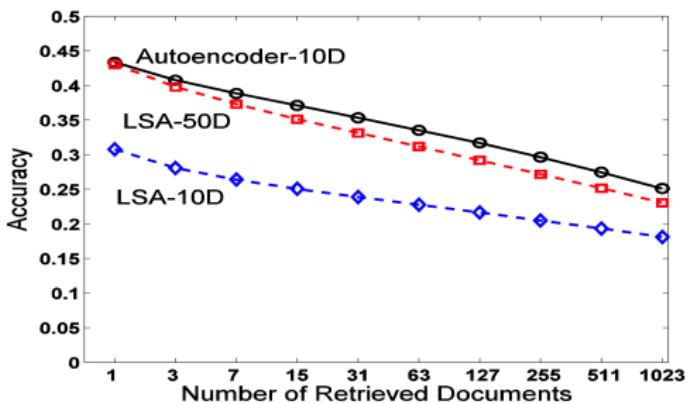
Autoencoder structure: 2000 — 500 — 250 — 125 — 2

<https://www.cs.toronto.edu/~hinton/science.pdf>

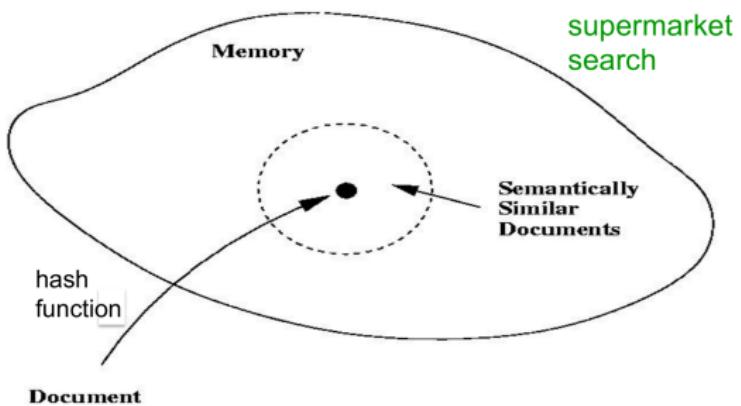
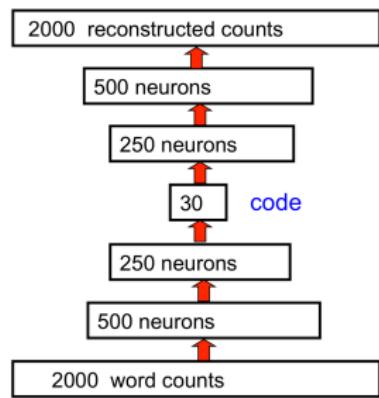
## Comparison with PCA



# Retrieval performance on 400,000 Reuters business news stories

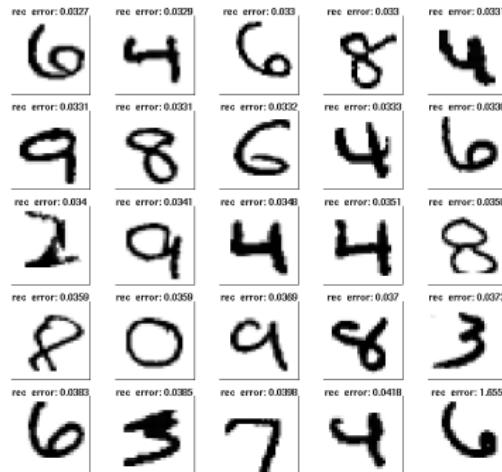


# Finding binary codes for documents



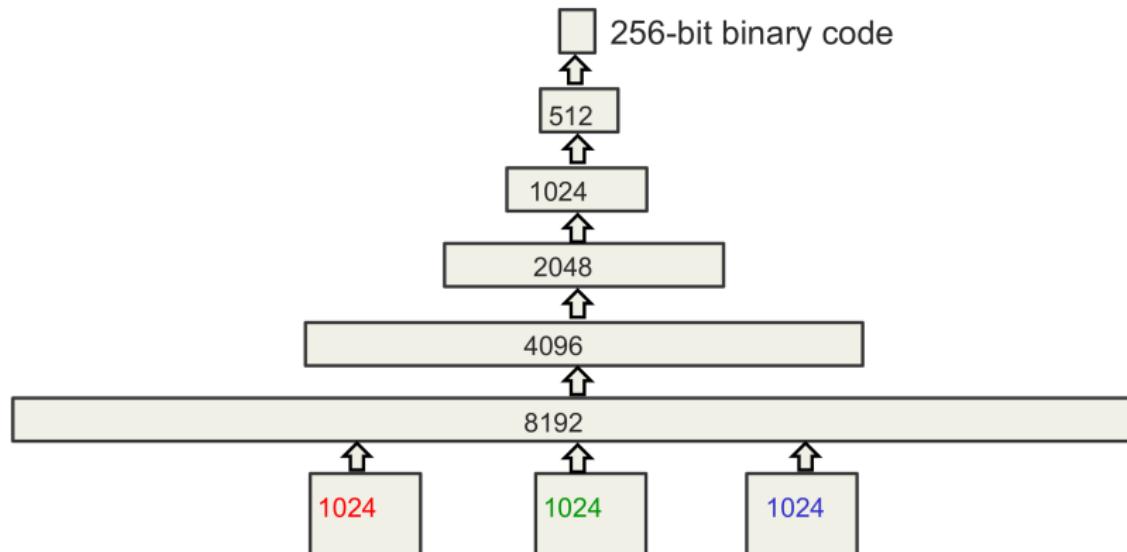
# Autoencoder: Why are they useful?

- ▶ unsupervised pretraining of weights (many unlabeled images, but only few labeled)
- ▶ anomaly detection

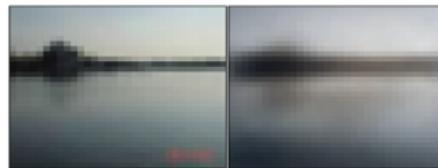
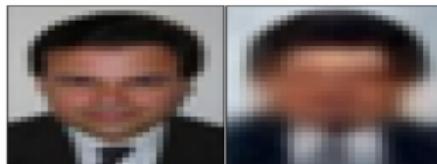


# Autoencoder: Why are they useful?

Image retrieval



## Reconstructions of 32x32 color images from 256-bit codes



retrieved using 256 bit codes



retrieved using Euclidean distance in pixel intensity space

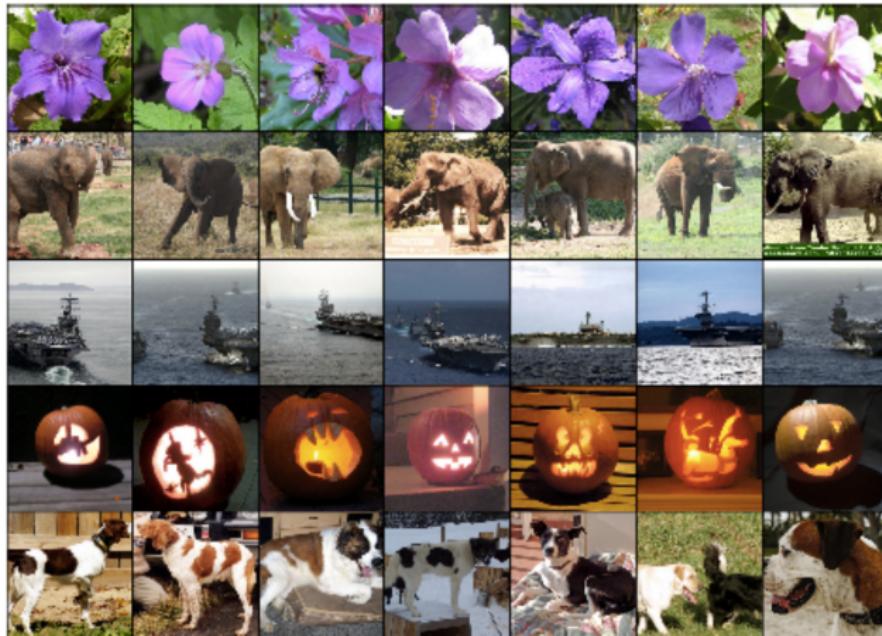


retrieved using 256 bit codes



retrieved using Euclidean distance in pixel intensity space



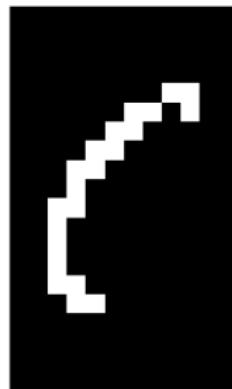
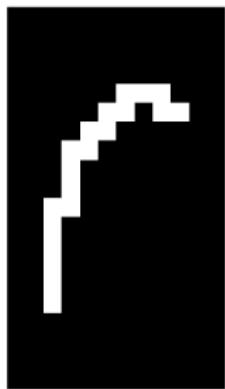


Leftmost column  
is the search  
image.

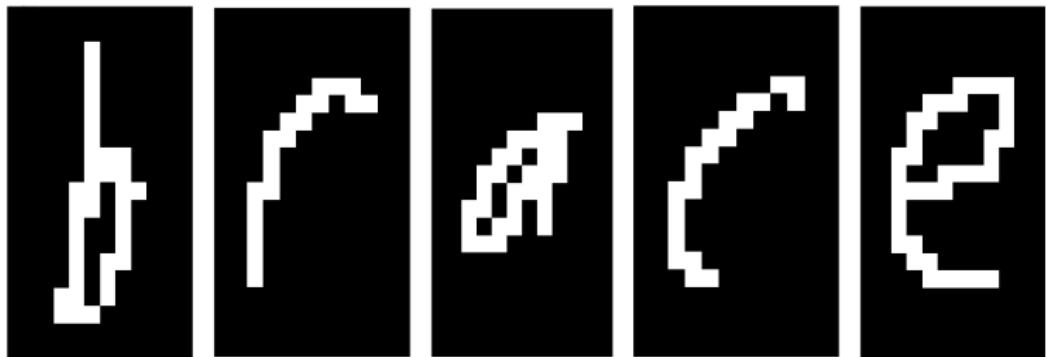
Other columns  
are the images  
that have the  
most similar  
feature activities  
in the last hidden  
layer.

## Sequence Models

## Handwritten character recognition

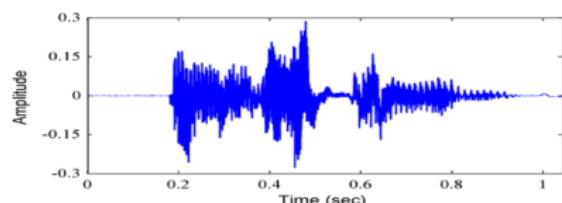


## Structured prediction



# Sequential data

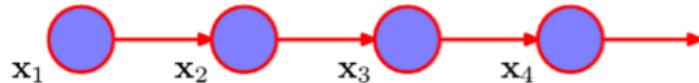
- ▶ time-series data (speech)
- ▶ characters in a sentence
- ▶ base pairs along a DNA strand



# Markov Model

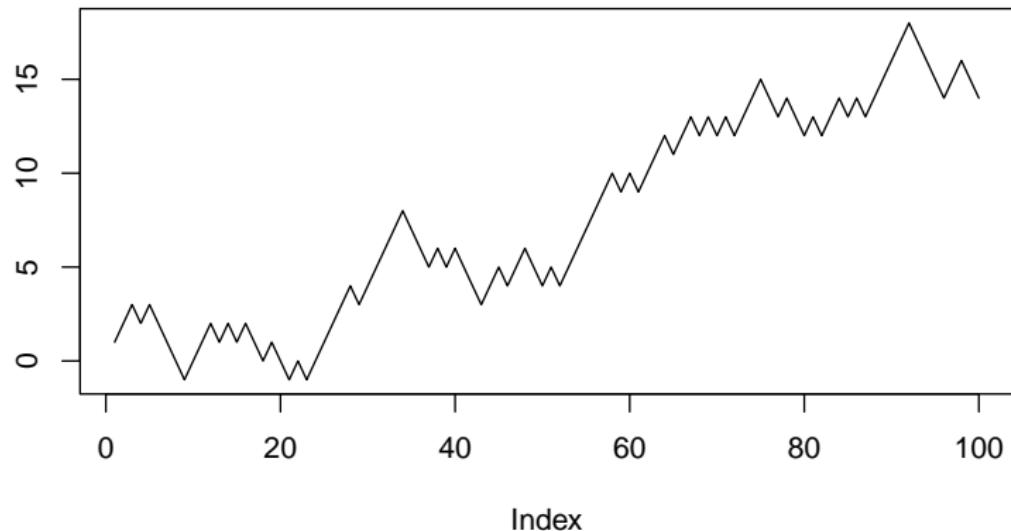
Joint distribution of n arbitrary random variables

$$\begin{aligned} P(X_1, X_2, \dots, X_n) &= P(X_1) \cdot P(X_2 | X_1) \cdot \dots \cdot P(X_n | X_{n-1}) \\ &= P(X_1) \cdot \prod_{i=2}^n P(X_i | X_{i-1}) \end{aligned}$$

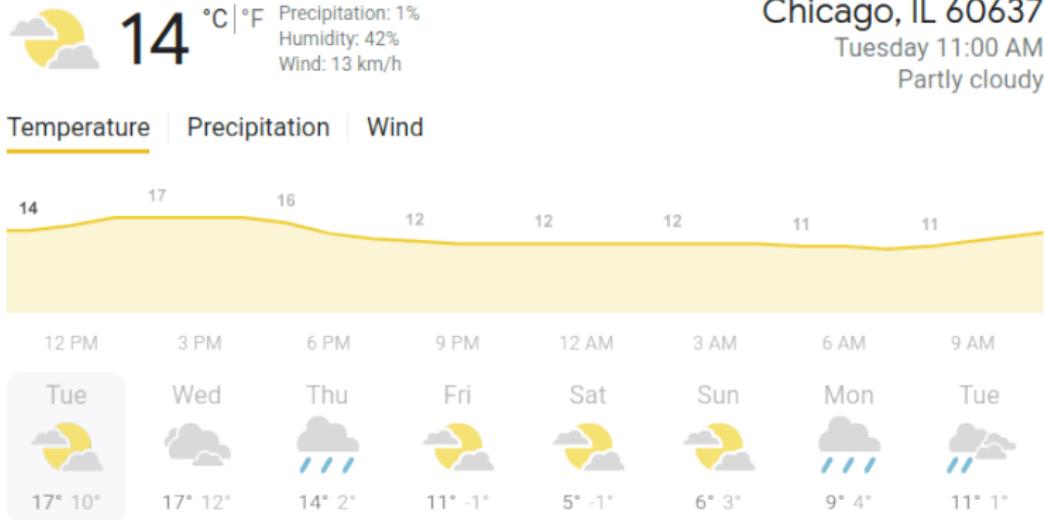


## Example

**Random walk model**



# Example



## Recurrent neural networks

# Recurrent neural networks

RNNs are very powerful, because they combine two properties:

- ▶ Distributed hidden state that allows them to store a lot of information about the past efficiently.
- ▶ Non-linear dynamics that allows them to update their hidden state in complicated ways.

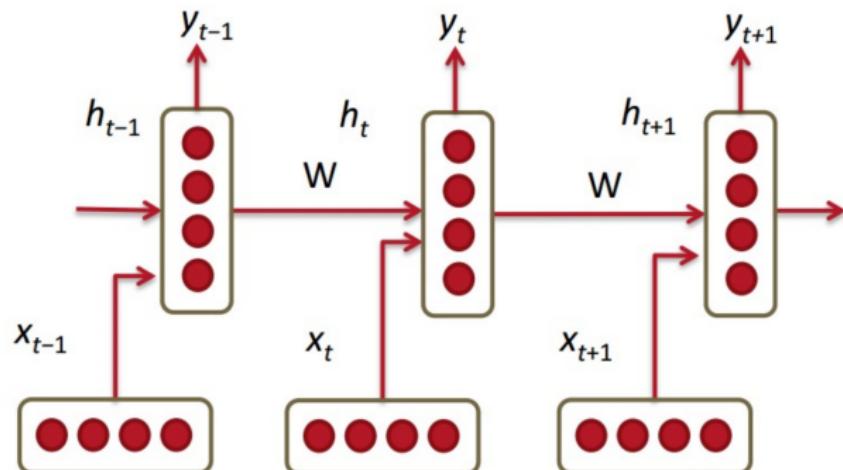
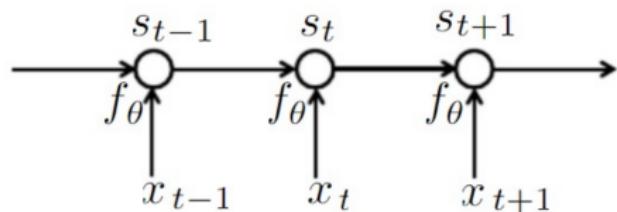


Figure: Richard Socher

# Recurrent neural networks as Dynamic systems

$$s_t = f_\theta(s_{t-1}, x_t)$$



The state contains information about the whole past sequence.

$$s_t = g_\theta(x_t, x_{t-1}, x_{t-s}, \dots, x_2, x_1)$$

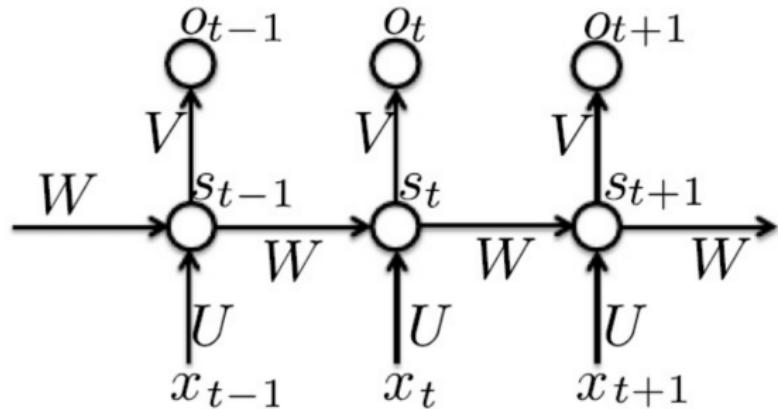
## Parameter sharing

We can think of  $s_t$  as a summary of the past sequence of inputs up to  $t$ .

If we define a different function  $g_t$  for each possible sequence length, we would not get any generalization.

If the same parameters are used for any sequence length allowing much better generalization properties.

# Recurrent Neural Networks



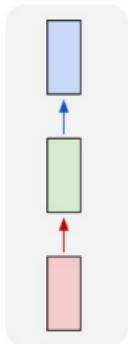
$$\mathbf{a}_t = \mathbf{b} + W\mathbf{s}_{t-1} + U\mathbf{x}_t$$

$$\mathbf{s}_t = \tanh(\mathbf{a}_t)$$

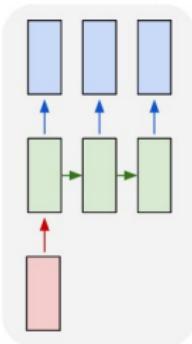
$$\mathbf{o}_t = \mathbf{c} + V\mathbf{s}_t$$

$$\mathbf{p}_t = softmax(\mathbf{o}_t)$$

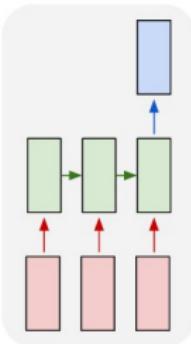
one to one



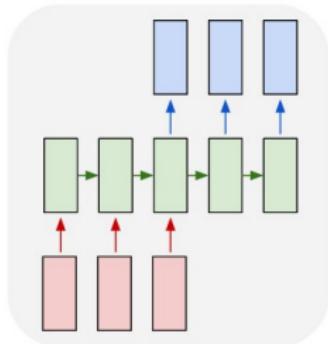
one to many



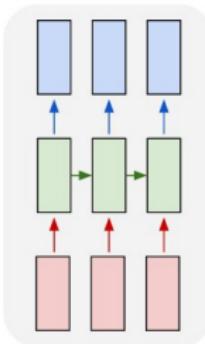
many to one



many to many



many to many



# Recurrent Neural Networks: Process Sequences

One to one

- ▶ vanilla neural network

One to many

- ▶ image captioning (image → sequence of words)

Many to one

- ▶ sentiment classification (sequence of words → sentiment)

Many to many

- ▶ machine translation (sequence of words → sequence of words)
- ▶ video classification on frame level

## Example: Sentiment Classification

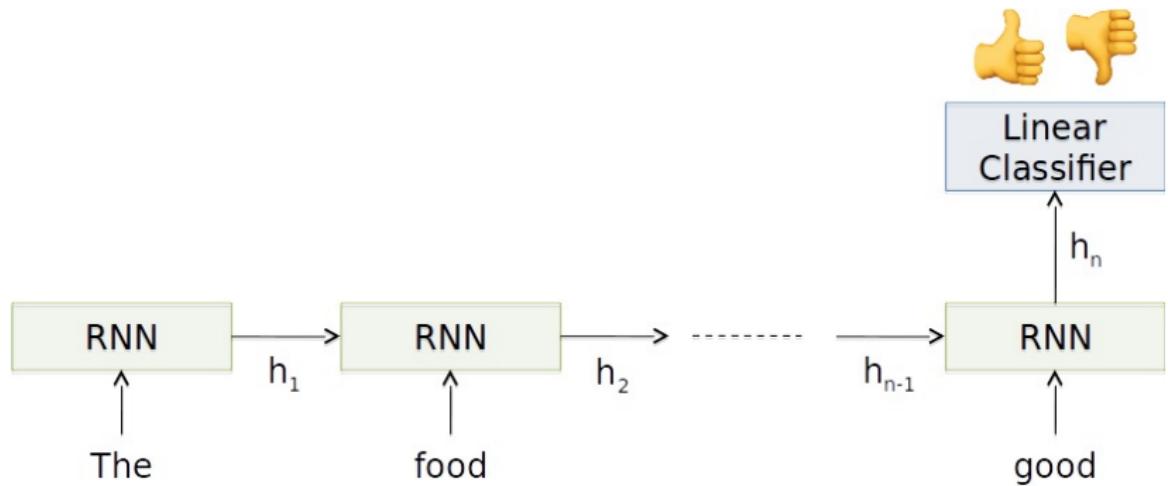
Classify a restaurant review from Yelp! OR movie review from IMDB as positive or negative

*Inputs:* Multiple words, one or more sentences

*Outputs:* Positive / Negative classification

“The food was really good”

## Example: Sentiment Classification



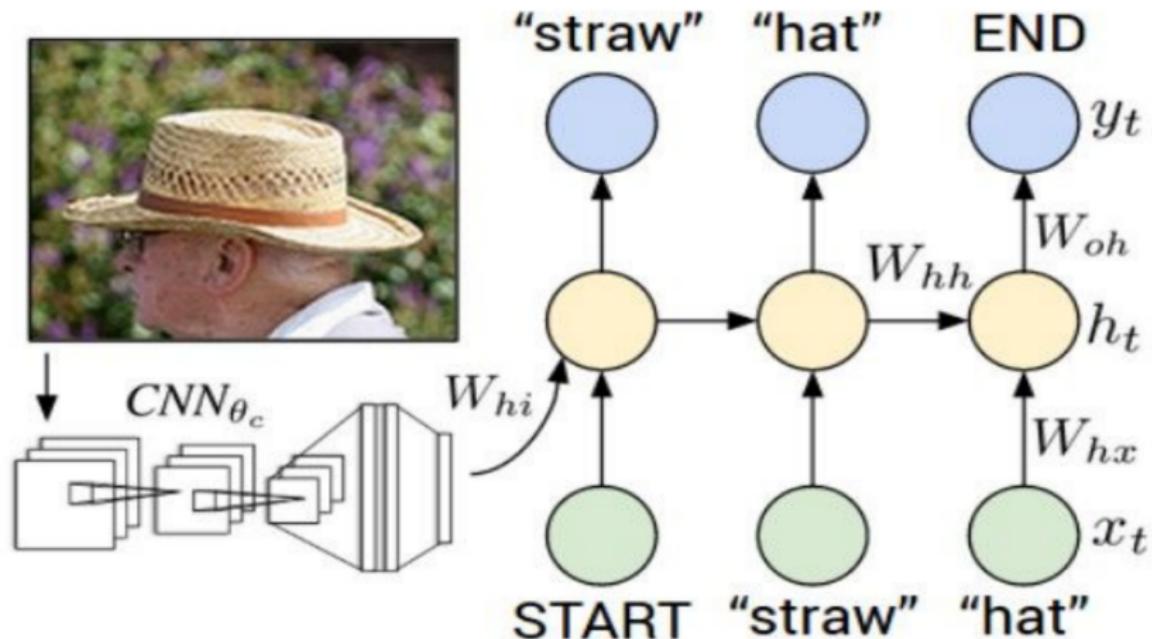
## Example: Image Captioning

Given an image, produce a sentence describing its contents

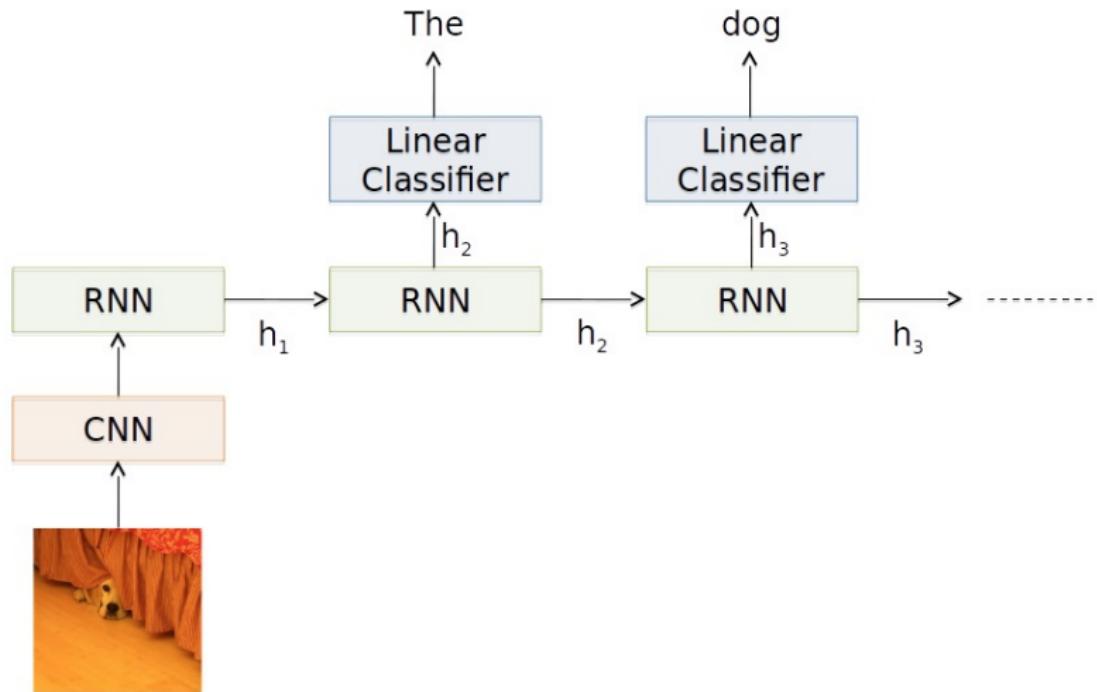
*Inputs:* Image feature (from a convolutional neural network)

*Outputs:* Multiple words (one sentence)

## Example: Image Captioning



## Example: Image Captioning



# Example: Image Captioning



A cat sitting on a suitcase on the floor



A cat is sitting on a tree branch



A dog is running in the grass with a frisbee



A white teddy bear sitting in the grass



Two people walking on the beach with surfboards



A tennis player in action on the court



Two giraffes standing in a grassy field



A man riding a dirt bike on a dirt track

# Example: Image Captioning



*A woman is holding a cat in her hand*



*A person holding a computer mouse on a desk*



*A woman standing on a beach holding a surfboard*



*A bird is perched on a tree branch*



*A man in a baseball uniform throwing a ball*

## Examples

Handwriting

<https://www.cs.toronto.edu/~graves/handwriting.html>

Image captioning

<https://cocodataset.org/#explore>